

Bayesian causal forests for estimating heterogeneous effects with joint treatments and interference

Abhi Jain

August 5, 2026

PhD Candidate

Department of Biostatistics

Boston University School of Public Health

Table of contents

1. Motivating problem
2. Causal inference with interference
3. Bayesian additive regression trees (BART)
4. Simulations
5. Case study: Effect of healthy food access on obesity rates in Texas
6. Discussion

Motivating problem

Effect of healthy food access on obesity rates

Poor healthy food access been associated with:

- high systolic blood pressure (Suarez et al., 2015)
- hypertension and cardiovascular disease (Kelli et al., 2019)
- obesity (Thomsen et al., 2016)

Challenges include:

- **Interference:** cross-border shopping
- **Heterogeneous effects:** magnitude and direction of effects can vary by community characteristics

Research question: What is the effect of healthy food retailers on ZIP Code obesity rates and how does it vary for different ZIP Code characteristics?

Causal inference with interference

Interference

- **No interference assumption:** The treatment/exposure of one unit does not impact the outcome of another unit
- Often violated in studies with spatial treatments or exposures




Figure 1: No interference




Figure 2: Spatial network interference

Notation and Potential outcomes

- $\mathbf{W}$ = adjacency matrix
 - w_{ij} = strength of connection between units i and j
 - if $w_{ij} = 0$, units i and j are unconnected
- $\mathcal{N}_i$ = set of units that are connected to unit i
 - Call this set the neighbors of unit i
- $(i, \mathcal{N}_i, \mathcal{N}_{-i})$ = partition of all units
- Potential outcome: $Y(A_i, \tilde{A}_i = f(\mathbf{A}_{\mathcal{N}_i}))$
 - $(A_i, \tilde{A}_i)$ = joint treatment
 - $f(\mathbf{A}_{\mathcal{N}_i})$ is a summarizing function of the neighboring exposures/treatments

- **Direct effect (DE)**: effect of unit-level treatment (A) on Y holding neighbor-level treatment fixed

$$DE(\tilde{A} = \tilde{a}) = E[Y(1, \tilde{a}) - Y(0, \tilde{a})]$$

- **Indirect (spillover) effect (IE)**: effect of neighbor-level treatment ($\tilde{A}$) holding unit-level treatment fixed

$$IE(A = a) = E[Y(a, 1) - Y(a, 0)]$$

- For simplicity, we consider binary A and $\tilde{A}$

Assumptions

1. Consistency: $Y_i = Y_i(\mathbf{A})$
2. **Neighborhood interference**: Given a function f , if $f(\mathbf{A}_{\mathcal{N}_i}) = f(\mathbf{A}'_{\mathcal{N}_i})$, then $Y(A_i, \mathbf{A}_{\mathcal{N}_i}) = Y(A_i, \mathbf{A}'_{\mathcal{N}_i})$ (Forastiere et al., 2021; Hudgens and Halloran, 2008)
3. Unconfoundedness of the joint treatment:
 $Y_i(a, \tilde{a}) \perp A_i, \tilde{A}_i | \mathbf{X} \quad \forall a, \tilde{a}, \quad \forall i$
4. Positivity: $f(a, \tilde{a} | \mathbf{X}) > 0 \quad \forall \mathbf{x} \quad \text{s.t.} \quad f_{\mathbf{X}}(\mathbf{x}) > 0$

Bayesian additive regression trees (BART)

- Ensemble machine learning method that models the outcome as a sum of weak learners
- Bayesian framework: places priors on tree structure and parameters and provides uncertainty quantification via credible intervals
- Often used for causal inference problems (Hill, 2011; Dorie et al., 2022)

Bayesian causal forests (BCF)

- Extension of BART designed for causal inference and heterogeneous treatment effects (Hahn et al., 2020)

$$Y_i = \mu(\mathbf{X}_i, \hat{\pi}(\mathbf{X}_i)) + \tau(\mathbf{X}_i)A_i + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2)$$

- $\mu(\mathbf{X})$ = prognostic function
- $\tau(\mathbf{X})$ = moderating function
- Include estimate of the propensity score in the prognostic function

Joint treatment Bayesian causal forests (jtBCF)

- Extension of Bayesian causal forests for joint treatments
- Separate moderating functions for both A , $\tilde{A}$, and their interaction

$$Y_i = \mu(\mathbf{X}_i, \hat{\pi}_A, \hat{\pi}_{\tilde{A}}) + \tau_1(\mathbf{X}_i)A_i + \tau_2(\mathbf{X}_i)\tilde{A} + \tau_3(\mathbf{X}_i)A\tilde{A} + \epsilon_i, \epsilon_i \sim N(0, \sigma^2)$$

- $\mu(\mathbf{X})$: prognostic effect
- $\tau_1(\mathbf{X})$: direct effect (own treatment)
- $\tau_2(\mathbf{X})$: indirect (spillover) effect
- $\tau_3(\mathbf{X})$: interaction effect

- We use the SoftBart R package (Linero, 2022) to implement our model
 - SoftBart R package also includes functionality to develop custom BART models
- Soft BART replaces binary splits with smooth gating functions (e.g., logistic functions) so each observation can partially belong to multiple leaves (Linero and Yang, 2018)

Simulations

Simulation setup

- Use Texas ZIP Codes as set of units ($n = 1927$)
- Simulate three ZIP Code-level covariates and compute three neighborhood-level covariates, which are neighborhood means of the ZIP Code-level covariates
- Consider linear/nonlinear $\mu(\mathbf{X})$
- Consider homogeneous/heterogeneous $\tau_1(X) - \tau_3(X)$
- Treatment:


$$A \sim \text{Bernoulli}(0.8 \Phi(-1.25 + \mu(X)/\text{sd}(\mu(X)) + 0.5X_5) + u_i/10)$$

$$\tilde{A}_i = \mathbf{1}(\bar{A}_{\mathcal{N}_i} > 0.5)$$


- Outcome:

$$y \sim N(\mu(X) + \tau_1(X)A + \tau_2(X)\tilde{A} + \tau_3(X)A\tilde{A}, 1)$$

Simulation results (Linear/Homogeneous case)



Simulation results (Linear/Heterogeneous case)







Coverage - Direct and Indirect effects

Setting	Estimand	Soft BART	Soft psBART	Soft jtBCF
Homo				
	$DE(\tilde{A} = 0)$	0.80	0.78	0.94
	$DE(\tilde{A} = 1)$	0.94	0.94	0.93
	$IE(A = 0)$	0.90	0.86	0.93
	$IE(A = 1)$	0.91	0.90	0.96
Hetero				
	$DE(\tilde{A} = 0)$	0.82	0.82	0.96
	$DE(\tilde{A} = 1)$	0.92	0.93	0.93
	$IE(A = 0)$	0.93	0.93	0.92
	$IE(A = 1)$	0.95	0.95	0.97





Case study: Effect of healthy food access on obesity rates in Texas

- **Unit of analysis:** Texas ZIP Codes in 2022
- **Outcome:** Obesity prevalence ($y \in (0, 100)$)
 - Model-based estimate adjusted for age, sex, race, and education
- **Exposures:**
 - A : Indicator for ≥ 1 healthy food retailer in the ZIP Code.
 - $\tilde{A}$: Indicator for ≥ 1 healthy food retailer in neighboring ZIP Codes.
- **Covariates:** Median household income, % SNAP, unemployment, % commuting by public transit, NDVI, % married, % veterans, and % with disability
- **Neighborhood:** Informed by driving times between ZIP Code centroids

Posterior distributions of direct/indirect effects



Results: Heterogeneous effects



Discussion

Conclusions

- jtBCF model offers bias reduction in estimating direct and indirect effects under interference compared to BART and psBART
- Provides strong coverage of direct and indirect effects
- Limitations:
 - Model has increased variability due to noise coming from three separate moderating functions
 - Motivating example has some weaknesses:
 - Data is ecological
 - Obesity rate at ZIP Code level is model-based estimate (takes into account race, sex, age, and education)
 - Food retailer list is not exhaustive; does not include farmer's markets or food banks

Thank you!



Shariq Mohammed



Sara Lodi



Wei Jin



QR code for slides

Any questions?

References

- Vincent Dorie, George Perrett, Jennifer L Hill, and Benjamin Goodrich. Stan and bart for causal inference: Estimating heterogeneous treatment effects using the power of stan and the flexibility of machine learning. *Entropy*, 24(12):1782, 2022.
- Laura Forastiere, Edoardo M Airoidi, and Fabrizia Mealli. Identification and estimation of treatment and interference effects in observational studies on networks. *Journal of the American Statistical Association*, 116(534):901–918, 2021.
- P Richard Hahn, Jared S Murray, and Carlos M Carvalho. Bayesian regression tree models for causal inference: Regularization, confounding, and heterogeneous effects (with discussion). *Bayesian Analysis*, 15(3):965–1056, 2020.
- Jennifer L Hill. Bayesian nonparametric modeling for causal inference. *Journal of Computational and Graphical Statistics*, 20(1):217–240, 2011.
- Michael G Hudgens and M Elizabeth Halloran. Toward causal inference with interference. *Journal of the american statistical association*, 103(482):832–842, 2008.

- Heval M Kelli, Jeong Hwan Kim, Ayman Samman Tahhan, Chang Liu, Yi-An Ko, Muhammad Hammadah, Samaah Sullivan, Pratik Sandesara, Ayman A Alkhoder, Fahad K Choudhary, et al. Living in food deserts and adverse cardiovascular outcomes in patients with cardiovascular disease. *Journal of the American Heart Association*, 8(4):e010694, 2019.
- Bora Kim. Estimating spillover effects in the presence of isolated nodes. *Spatial Economic Analysis*, pages 1–15, 2025.
- Antonio R Linero. Softbart: Soft bayesian additive regression trees. *arXiv preprint arXiv:2210.16375*, 2022.
- Antonio R Linero and Yun Yang. Bayesian regression tree ensembles that adapt to smoothness and sparsity. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 80(5):1087–1110, 2018.
- Jonathan J Suarez, Tamara Isakova, Cheryl AM Anderson, L Ebony Boulware, Myles Wolf, and Julia J Scialla. Food access, chronic kidney disease, and hypertension in the us. *American journal of preventive medicine*, 49(6):912–920, 2015.
- Michael R Thomsen, Rodolfo M Nayga Jr, Pedro A Alviola IV, and Heather L Rouse. The effect of food deserts on the body mass index of elementary schoolchildren. *American Journal of Agricultural Economics*, 98(1):1–18, 2016.

$$\mu(\mathbf{x}) = \begin{cases} 1 + 0.7X_1 - 0.6X_2 + X_3 + 0.3X_4, & \text{linear} \\ 0.5 + 0.6|X_1 + 1| - 0.8\sin(\pi X_2) + X_3 - 0.6X_4 + 0.3X_1X_4, & \text{nonlinear} \end{cases}$$

$$\tau(\mathbf{x}) = \begin{cases} 0.4A + 0.4\tilde{A} + 0.2A\tilde{A}, & \text{homo} \\ 0.3(1 + \frac{X_1}{2} + X_3)A + 0.3(1 + \frac{X_2}{2} + X_3)\tilde{A} + 0.15(X_3 + X_6)A\tilde{A}, & \text{hetero} \end{cases}$$

Results: Overall effects

- $DE(\tilde{A} = 0) = -0.17(-0.84, 0.54)$
- $DE(\tilde{A} = 1) = -0.56(-1.03, -0.10)$
- $IE(A = 0) = -0.16(-0.86, 0.50)$
- $IE(A = 1) = -0.56(-1.25, 0.17)$

Isolated nodes

- Since we use driving times to inform the neighborhood adjacency, some ZIP Codes are isolated (no neighbors)
- Often it is assumed $\tilde{A} = 0$ for isolated nodes
- We drop these ZIP Codes from the analysis since it has been shown that estimation of spillover effects is biased when we assume $\tilde{A} = 0$ for isolated nodes (Kim, 2025)